

Floating Point Numbers Representation

Philippe Bonnet, bonnet@di.ku.dk
HPPS 2025 – 2a

(with slides from T.Henriksen)

Outline

- Numerical errors
- The fixed-point dilemma
- Floating-point representation
- IEEE754 and beyond

Numerical errors with floats

- Comparison errors
 - Never check for floating point numbers equality
 - Check that their difference is within a given bound
- Loop errors
 - Accumulation of errors in a loop
- Small number may get lost in relation to large number

See examples in 2a-code/

Fractional numbers

$$123.456 = 1 \cdot 10^2 + 2 \cdot 10^1 + 3 \cdot 10^0 + 4 \cdot 10^{-1} + 5 \cdot 10^{-2} + 6 \cdot 10^{-3}$$

Generally

$$a_{m-1} \cdots a_0.a_{-1} \cdots a_{-n} = \sum_{i=-n}^{m-1} a_i \cdot 10^i$$

Even more generally, for radix r

$$a_{m-1} \cdots a_0.a_{-1} \cdots a_{-n} = \sum_{i=-n}^{m-1} a_i \cdot r^i$$

Fractional binary numbers

Weight	2^{m-1}	2^{m-2}	...	4	2	1	$1/2$	$1/4$	$1/8$...	2^{-n}
Digits	b_{m-1}	b_{m-2}	...	b_2	b_1	b_0	b_{-1}	b_{-2}	b_{-3}	...	b_{-n}

Representation

- Bits to the right of “binary point” represents fractional powers of 2.
- Represents rational number

$$\underbrace{b_{m-1} \cdots b_0}_{\text{integral part}} \cdot \underbrace{b_{-1} \cdots b_{-n}}_{\text{fraction part}} = \sum_{i=-n}^{m-1} b_i \cdot 2^i$$

Fractional binary numbers

Value	Representation
$5\frac{3}{4}$	101.11_2

$2\frac{7}{8}$	10.111_2
----------------	------------

$1\frac{7}{16}$	1.0111_2
-----------------	------------

Observations

- Divide by 2 by logical shifting right.
- Multiply by 2 by shifting left.
- Numbers of form $0.111\dots$ are just below 1.0.
 - ▶ $1/2 + 1/4 + 1/8 + \dots + 1/2^n + \dots \sim 1.0$.
 - ▶ Use notation $1.0 - \epsilon$.

Fractional binary numbers -- limitations

Limitation #1

- Can only represent fractional part of form $x/2^k$
- Other rational numbers have repeating binary representation

Problem for
fixed size
Types!

Value	Representation
$\frac{1}{3}$	0.0101010101[01] \cdots_2
$\frac{1}{5}$	0.001100110011[0011] \cdots_2
$\frac{1}{10}$	0.0001100110011[0011] \cdots_2

Limitation #2

- Just one setting of binary point within the w bits.
 - ▶ Limited range of numbers—very small values? Very large?

The fixed-point dilemma

Consider $w = 8$

1 bit for fraction

- Largest number: $1111111.1_2 = 127.5_{10}$
- Increment: $0000000.1_2 = 0.5_{10}$

7 bits for fraction

- Largest number: $1.1111111_2 = 1.9921875_{10}$
- Increment: $0.0000001_2 = 0.0078125_{10}$

4 bits for fraction

- Largest number: $1111.1111_2 = 15.9375_{10}$
- Increment: $0000.0001_2 = 0.0625_{10}$

Uniform increments (i.e., same precision) everywhere
=> Relative precision is worse close to 0!

Floating-point representation

Numerical form

Significand is
also called
Mantissa.

$$(-1)^s \cdot m \cdot 2^e$$

- **Sign bit** s determines whether number negative or positive.
- **Significand** m normally a fractional value in range $[1, 2)$.
- **Exponent** e weights value by power of two.

Encoding

- Most significant bit is sign bit.
- E field encodes e (but is not equal to e).
- T field encodes m (but is not equal to m).

Numerical
errors!


How are E and
 T encoded?




$$(-1)^s \cdot m \cdot 2^e$$

Precision (IEEE Standard 754)

32-bit single precision: float



64-bit double precision: double






$$(-1)^s \cdot m \cdot 2^e$$



Floating-point representation - NaN

← very positive e very negative $e \rightarrow$ ← very negative e very positive $e \rightarrow$



- Not a Number (NaN)




- $E = <111 \dots 1>$
- $T \neq <000\dots 0>$
- NaN is for the cases where no numerical value is defined, e.g., $0^{*\infty}$, $\sqrt{-1}$
- NaN is not greater than or lower than any number. Sign bit is irrelevant.

$$(-1)^s \cdot m \cdot 2^e$$



Floating-point representation - Infinity

← very positive e very negative $e \rightarrow$ ← very negative e very positive $e \rightarrow$



- **Infinity ($+\infty, -\infty$)**




- $E = <111 \dots 1>$
- $T = <000\dots 0>$
- $\frac{1}{0} = \infty$; $\frac{-1}{0} = -\infty$

$$(-1)^s \cdot m \cdot 2^e$$



Floating-point representation - Zero

← very positive e very negative $e \rightarrow$ ← very negative e very positive $e \rightarrow$



- Zero




- $E = <000\dots 0>$
- $T = <000\dots 0>$
- Note that -0 and $+0$ have distinct encodings!
At bit-level, $-0 \neq 0$ but in C: $0 == -0$

$$(-1)^s \cdot m \cdot 2^e$$



Floating-point representation – Normal values

← very positive e very negative $e \rightarrow$ ← very negative e very positive $e \rightarrow$






- E encoded as biased value: $e = B2U(E) - b$
 - See *last week's slides*
 - Single precision: $b = 127 \Rightarrow E \in [-126; 127]$
 - Double precision: $b = 1023 \Rightarrow E \in [-1022; 1023]$
- T encoded with extra leading 1: $m = 1 + B2U(T) \cdot 2^{1-p}$
 - $p = 1 + \text{sizeof}(T)$ - single precision $p = 1 + 23 = 24$; $p = 1 + 52 = 53$
 - $T = <000\dots 0> \Rightarrow m = 1$; $T = <111\dots 1> \Rightarrow m = 2 - \varepsilon$

$$(-1)^s \cdot m \cdot 2^e$$




Floating-point representation – Subnorm

\leftarrow very positive e very negative $e \rightarrow$ \leftarrow very negative e very positive $e \rightarrow$



- $E = <000..0>$: $e = 1 - b$
- T encoded with extra leading 0: $m = B2U(T) \cdot 2^{1-p}$
 - $T = <000..0> \Rightarrow 0$
 - $T \neq <000..0> \Rightarrow$ Numbers close to 0

FP8 E4M3



8-bit floating-point representation


- Sign bit is the most significant bit (leftmost).
- The next four bits are E with a bias of 7.
- The last three bits are T .

Alternative: E5M2 – E encoded on 5 bits, T on 2 bits.

FP8 values


	<i>S</i>	<i>E</i>	<i>T</i>	<i>e</i>	Value	
Denormalised	0	0000	000	-6	0	
	0	0000	001	-6	$1/8 \cdot 1/64 = 1/512$	closest to zero
	0	0000	010	-6	$2/8 \cdot 1/64 = 2/512$	
	...					
Normalised	0	0000	111	-6	$7/8 \cdot 1/64 = 7/512$	largest denorm
	0	0001	000	-6	$8/8 \cdot 1/64 = 8/512$	smallest norm
	0	0001	001	-6	$9/8 \cdot 1/64 = 9/512$	
	...					
	0	0110	110	-1	$14/8 \cdot 1/2 = 14/16$	
	0	0110	111	-1	$15/8 \cdot 1/2 = 15/16$	Closest to 1
	0	0111	000	0	$8/8 \cdot 1 = 1$	
	0	0111	001	0	$9/8 \cdot 9/8 = 1$	Closest to 1
	0	0111	010	0	$10/8 \cdot 10/8 = 1$	
	...					
	0	1110	110	7	$14/8 \cdot 128 = 224$	
	0	1110	111	7	$15/8 \cdot 128 = 240$	
	0	1111	000	N/A		∞

FP8 – non-uniform distribution of values



Example:

- 150 represented as 144
- 151 represented as 144
- 152 represented as 144
- 153 represented as 160



Take-aways

- The fixed representation of decimal number leads to numerical errors
- Floating-points are implemented in most languages
- Floating-points can represent infinity and NaN
- Floating points have a dynamic range
- Small floating-point types used for machine learning
 - Compact representation vs. loss of precision