

Program Representation


Philippe Bonnet, bonnet@di.ku.dk
HPPS 2025 – 3a

Outline

- Modern processors
- Assembly
- Compilers & Interpreters
- Experiments

Instruction Set Architecture


- Assembly Language View
 - Processor state
 - Registers, memory, ...
 - Instructions
 - addq, pushq, ret, ...
 - How instructions are encoded as bytes
- Layer of Abstraction
 - Above: how to program machine
 - Processor executes instructions in a sequence
 - Below: what needs to be built
 - Use variety of tricks to make it run fast
 - E.g., execute multiple instructions simultaneously




<https://www.tele-task.de/lecture/video/8963/>

Modern processors


Chiplet-based architecture



(a) AMD EPYC Milan



(b) Intel Sapphire Rapids



(c) ARM Graviton 3


Parallelism

- Multiple cores => multiple parallel **threads** of executions
- Out of order execution, speculation, multi-issue => multiple parallel **instructions** (at the cost of non-determinism)
- Single-instruction multiple-data (SIMD) => multiple parallel operations on **data**

Locality:

- **Cache hierarchy** per core and across cores within a processor

Processor as Interpreter



Processor abstraction

Programmer-Visible State

PC: Program counter


- Address of next instruction, 8B
- Called “RIP” (x86-64)

Registers

- Data (to be) processed (arithmetic, logic, control)
- Each register contains 8B

Condition codes

- Store status information about most recent arithmetic operation
- Used for conditional branching




<https://en.wikichip.org/wiki/x86>

Intel x86-64

Little endian
variable length
instruction set

Binary representation



Textual representation

Assembly language

For historical reasons, r0-r7 are called original registers.

They have the following names:

- ax: register a
- bx: register b
- cx: register c
- dx: register d
- bp: register base pointer (start of stack)
- sp: register stack pointer (current location in stack, grow downwards)
- si: register source index (source for data copies)
- di: register destination index (destination for data copies)

Register values can be accessed as:

- **8B:**
 - original registers: prefix r rax, rsp, rsi
 - Other registers: no suffix r8, r15
- **4B:**
 - original registers: prefix e eax, esp, esi
 - other registers: suffix d r8d, r15d
- **2B:**
 - original registers: No prefix ax, sp, si
 - other registers: suffix w r8w, r15w
- **1B (high byte):**
 - original registers (bits 8-15 from ax-dx) ah, bh, ch, dh
- **1B (low byte):**
 - original registers (bits 0-7 from ax-dx) al, bl, cl, dl
 - other registers: suffix b r8b, r15b

The screenshot shows the Compiler Explorer interface. On the left, the C source code for a square function is displayed:1 /* Type your code here, or load an example. */
2 int square(int num) {
3 return num * num;
4 }A large blue downward-pointing arrow is positioned between the 'Binary representation' and 'Textual representation' sections. In the center, a box contains information about AT&T syntax:

The GNU tools (gcc, gdb) use AT&T Syntax for assembly
`movq %rsp, %rbp`

It is of the form
OPERATOR source, destination

Register names are prefixed with %

On the right, the assembly output for the x86-64 gcc 15.2 compiler is shown:1 square:
2 pushq %rbp
3 movq %rsp, %rbp
4 movl %edi, -4(%rbp)
5 movl -4(%rbp), %eax
6 imull %eax, %eax
7 popq %rbp
8 retAt the bottom, the Godbolt URL godbolt.org is visible.

Assembly language

Three classes of instructions:


1. Transfer between memory and register
 - Load/store data: register <-> memory
 - Push/pop: register <-> **stack**
2. Arithmetic and comparison functions
3. Transfer control
 - Jumps to/from procedures
 - Conditional branches

How are registers organized?
How is memory addressed?

How are procedure calls
organized?

Virtual memory

- Every process in Linux has a similar view of **virtual memory**.
- Assembly manipulates **virtual addresses** in load/store.
- Within the CPU, a **memory management unit** translates virtual addresses into physical addresses.



Structured programming


Functions transform inputs into output

- ⇒ **Memory allocation** for input and output parameters
Memory allocation for local variables (defined within the function)
- ⇒ **Control flow** (which instruction to execute next)
when calling a function
when returning from a function

These two problems are solved with the **stack**

The Stack

- Region of memory managed with stack discipline (LIFO)
- Grows towards lower addresses
- Register %rbp points to the bottom of the stack
- Register %rsp points to the top of the stack, it is the **stack pointer**



Stack - Push

Operator push

- `pushq %r10`

1. Decrement `%rsp` by 8
2. Write contents of `%r10` at address given by `%rsp`

Stack Pointer: `%rsp`



Stack “Bottom”



Stack “Top”


Stack - Pop

Operator pop

- `popq %r10`

1. Copy contents at address `%rsp` to `%r10`
2. Increment `%rsp` by 8

Stack Pointer: `%rsp`



Stack “Bottom”



Stack “Top”


Example

```
yoo (...)  
{  
    •  
    •  
    who () ;  
    •  
    •  
}
```

```
who (...)  
{  
    • • •  
    amI () ;  
    • • •  
    amI () ;  
    • • •  
}
```

```
amI (...)  
{  
    •  
    •  
    amI () ;  
    •  
    •  
}
```



Example
Call Chain





Procedure **amI ()** is recursive

Example


```
yoo (...)  
{  
    :  
    :  
    who ();  
    :  
    :  
}
```





Example




Stack





Example




Stack





Example





Stack




Example





Example




Stack





Example




Stack





Example




Stack





Example




Stack



Example



Stack





Procedure Call

- Instructions:
 - **call**: push **return address** on stack; jump to label/address
Return address is address of instruction right after call instruction
 - **ret**: pop return address from stack into instruction pointer

Procedure Call Example

```
804854e: e8 3d 06 00 00    call   8048b90 <main>
8048553: 50                  pushl  %eax
```

call 8048b90



%rsp 0x108


%rsp 0x100

%rip 0x80854e

%rip 0x8048b90

Procedure Call Example

8048591:	c3	ret
----------	----	-----



Calling a function (x86-64)

To call a function:

1. **Place the first six integer or pointer parameters in %rdi, %rsi, %rdx, %rcx, %r8 and %r9**
2. Push onto the stack subsequent parameters and parameters larger than 8B (in order).
3. Execute the call instruction, which:
 - Pushes the return address onto the stack
 - Jumps to the start of the specified function (i.e., update instruction pointer)

Executing a function

- The C run-time system introduces instruction to set-up and clean-up the stack in each procedure.
- Set-up consists in allocation and initialization of a stack-frame. Clean-up consists in deallocating a stack frame.
- A stack-frame is the space needed on the stack by a procedure for storing:
 - The return address
 - (some) parameters
 - Local variables


Stack Frame

Caller:

- Arguments
 - pushed by program (if needed)
- Return address
 - pushed by call

Callee:

- Previous frame pointer (%rbp)
- Other callee-save registers (%rbx, %r12-15)
- Space for local variables
- Arguments for next function (when about to call another function)



COMPILER EXPLORER Add... More Templates

Sponsors intel. Cppcon Google Share Policies Other

C source #1

```
1 /* Type your code here, or load an example. */
2 int square(int num) {
3     return num * num;
4 }
5
6 int main(void) {
7     int x = 4;
8     return square(x);
9 }
10
11
12
```

x86-64 gcc 15.2 (Editor #1)

x86-64 gcc 15.2 Compiler options...

```
1 square:
2     pushq %rbp
3     movq %rsp, %rbp
4     movl %edi, -4(%rbp)
5     movl -4(%rbp), %eax
6     imull %eax, %eax
7     popq %rbp
8     ret
9 main:
10    pushq %rbp
11    movq %rsp, %rbp
12    subq $16, %rsp
13    movl $4, -4(%rbp)
14    movl -4(%rbp), %eax
15    movl %eax, %edi
16    call square
17    leave
18    ret
```

Output (0/0) x86-64 gcc 15.2 i - cached (4040B) ~253 lines filtered

godbolt.org

Compilers and interpreters

- Compiled languages

- + Almost always faster.
- Require compilation after every change.
- Usually cannot run program fragments in isolation.
- Tend to have more restrictions (e.g. static typing).
- Much more difficult to implement.

- Interpreted languages

- Usually slow.
- + Can run immediately.
- + Can easily run fragments (e.g. single functions) in isolation.
- + Much easier to implement.

(source T.Henriksen)

The Collatz conjecture

$$f(n) = \begin{cases} \frac{n}{2} & \text{if } n \text{ is even} \\ 3n + 1 & \text{if } n \text{ is odd} \end{cases}$$

Conjecture: if we apply this function to some number greater than 1, we will eventually reach 1.

Listing 1: collatz.py

```
import sys

def collatz(n):
    i = 0
    while n != 1:
        if n % 2 == 0:
            n = n // 2
        else:
            n = 3 * n + 1
        i = i + 1
    return i

k = int(sys.argv[1])
for n in range(1, k):
    print(n, collatz(n))
```

Listing 2: collatz.c

```
#include <stdio.h>
#include <stdlib.h>

int collatz(int n) {
    int i = 0;
    while (n != 1) {
        if (n % 2 == 0) {
            n = n / 2;
        } else {
            n = 3 * n + 1;
        }
        i++;
    }
    return i;
}

int main(int argc, char** argv) {
    int k = atoi(argv[1]);
    for (int n = 1; n < k; n++) {
        printf("%d %d\n", n, collatz(n));
    }
}
```

```

htl719@kumac ~/D/C/H/h/w/lec-tue> make python-run
time python3 ./collatz.py 100000 >/dev/null

real    0m1.268s
user    0m1.247s
sys     0m0.009s
htl719@kumac ~/D/C/H/h/w/lec-tue> make run
time ./collatz 100000 >/dev/null

real    0m0.041s
user    0m0.036s
sys     0m0.001s
htl719@kumac ~/D/C/H/h/w/lec-tue> make optim-run
time ./o-collatz 100000 >/dev/null

real    0m0.015s
user    0m0.012s
sys     0m0.001s

```

$$\text{Speedup} = \frac{1.268}{0.015} = 84.5$$

```

plain: collatz.c
      gcc collatz.c -o collatz

optim: collatz.c
      gcc -O3 -ffast-math collatz.c -o o-collatz

```

1st run

```

htl719@kumac ~/D/C/H/h/w/lec-tue> make python-run
time python3 ./collatz.py 100000 >/dev/null

real    0m1.449s
user    0m1.289s
sys     0m0.020s
htl719@kumac ~/D/C/H/h/w/lec-tue> make run
time ./collatz 100000 >/dev/null

real    0m0.224s
user    0m0.041s
sys     0m0.004s
htl719@kumac ~/D/C/H/h/w/lec-tue> make optim-run
time ./o-collatz 100000 >/dev/null

real    0m0.144s
user    0m0.014s
sys     0m0.004s

```

Mixing C and Python

Compiling C programs into object file that can be linked statically (at compile time)

or

shared object file that can be linked from Python (at run-time)

```
object: collatz.c
    gcc -c collatz.c -o collatz.o

reloc: collatz.c
    gcc -fPIC -shared collatz.c -o libcollatz.so
```

Linking libraries from Python

```
1 import sys
2 import ctypes
3
4 c_lib = ctypes.CDLL('./libcollatz.so')
5
6 if __name__ == '__main__':
7     k = int(sys.argv[1])
8     sum = 0
9     for i in range(1, k):
10         sum += c_lib.collatz(i)
11     print(sum)
~
```

Mixing Python and C


```
ntl719@kumac ~/D/C/H/h/w/lec-tue> time python3 collatz-ffi.py 100000
10753712
python3 collatz-ffi.py 100000  0.07s user 0.01s system 22% cpu 0.374 total
```

- x5 slower compared to C
- X20 faster compared to Python
 - Faster if we made fewer “foreign” calls, but each took more time.
 - Ideal case is single foreign function call that operates on many values.
 - **This is exactly how NumPy works!**

Experiments

Experimental framework

- System
 - Parameters (fixed)
 - Factors (take different values in an experiment)
- Workload
- Metric
 - Time / space
 - Throughput / latency
- Experiments
 - Studying the impact of modifying **one** factor at a time



(source: gemini)

Experimental results:

1. What you expect
2. What you observe
3. How you explain the discrepancies